

Foundations of Operations Research

Lorenzo Rossi and everyone who kindly helped!

2022/2023

Last update: 2022-10-17

These notes are distributed under Creative Commons 4.0 license - CC BY-NC 



no alpaca has been harmed while writing these notes

Contents

1	Introduction	1
1.1	Algorithm	1
1.2	Dynamic Programming	1
1.3	Complexity of algorithms	1
1.3.1	Big-O notation	1
1.4	Complexity classes	2
2	Graph and Network Optimization	3
2.1	Graphs	3
2.1.1	Graphs representation	4
2.1.2	Graph reachability problem	4
2.1.2.1	Complexity analysis	6
2.2	Subgraphs and Trees	6
2.3	Properties of trees	6
2.4	Optimal cost spanning tree	9
2.4.1	Prim's algorithm	9
2.4.1.1	Correctness of Prim's algorithm	9
2.4.1.2	Implementation in quadratic time	10
2.4.2	Optimality condition	10
2.5	Optimal paths	12
2.5.1	Dijkstra's algorithm	12
2.5.1.1	Correctness of Dijkstra's algorithm	13
2.5.1.2	Example of Dijkstra's algorithm	14
2.5.2	Floyd-Warshall's algorithm	16
2.5.2.1	Correctness of Floyd-Warshall's algorithm	16
2.6	Optimal paths in directed, acyclic graphs	16
2.6.1	Topological ordering method	18
2.6.2	Dynamic programming for shortest path in <i>DAGs</i>	18
2.6.2.1	Optimality of the algorithm	19
2.7	Project planning	19
2.7.1	Optimal paths	19
2.7.1.1	Critical path method - <i>CPM</i>	20
2.7.1.2	Critical paths	21
2.7.1.3	Gantt charts	21
3	Linear Programming	22
3.1	Optimization problems	22
3.2	Assumptions of <i>LP</i> models	23
3.3	Equivalent Forms	23
3.3.1	Transformation rules	23
3.4	Graphical solutions	24
3.4.1	<i>Weyl-Minkowski</i> Theorem	24
3.4.2	Geometry of <i>LP</i>	26
3.4.3	Four types of <i>LP</i>	26
3.4.4	Basic feasible solutions and polyhedra vertices	28
3.4.5	Algebraic characterization of vertices	28
3.4.6	Basic solutions	28
3.5	Simplex method	29
3.5.1	Optimality test	29
3.5.2	General case	30
3.5.3	Changes of basis	31
3.5.4	Pivoting operation	31
3.5.5	Moving to an adjacent vertex	32
3.5.6	Tableau representation	32

3.5.7	The simplex algorithm	32
3.5.8	Degenerate basic feasible solutions and convergence	33
3.5.8.1	Anti cycling rule	34
3.5.9	Two phase simplex algorithm	34
3.6	Network flows	34
3.6.1	Linear programming model	35

1 Introduction

1.1 Algorithm

An algorithm for a problem is a sequence of instructions that allows to solve any of its instances. The execution time of an algorithm depends on various factors, most notably the instance and the computer.

Properties:

- An algorithm is **exact** if it provides an optimal solution for every instance.
 - otherwise is **heuristic**
- A **greedy algorithm** constructs a feasible solution iteratively, by making at each step a *locally optimal* choice, without reconsidering previous choices
 - for most *discrete optimization problems*, greedy type algorithms yield a feasible solution with no guarantee of optimality

1.2 Dynamic Programming

Proposed by *Richard Bellman* in 1950, **dynamic programming** (or *DP*) is a method for solving optimization problems, composed of a sequence of decisions, by solving a set of recursive equations.

DP is applicable to any sequential decision problem, for which the optimality property is satisfied; as such, it has a wide range of applications, including scheduling, transportation, and assignment problems.

1.3 Complexity of algorithms

In order to analyze an algorithm, it's necessary to consider its complexity as a function of the size of the instance (*the size of the input*), independently of the computer; the complexity is defined as the number of elementary operations by assuming that each elementary operation takes a constant time.

Since it's hard to determine the exact number of elementary operations, an additional assumption is made: only the asymptotic number of elementary operations in the worst case (for the worst instances) is considered. The complexity evaluation is then performed by looking for the function $f(n)$ that best approximates the upper bound on the number of elementary operations n for the worst instances.

1.3.1 Big-O notation

A function f is of order of g , written $f(n) = \mathcal{O}(g(n))$ if there exists a constant $c > 0$ and a constant $n_0 > 0$ such that $f(n) \leq c \cdot g(n)$ for all $n \geq n_0$.

An illustration of the big-O notation is shown in Figure 1.




Figure 1: Big-O notation

n	n^2	2^n
1	$1\ \mu s$	$1\ \mu s$
10	$100\ \mu s$	$1.024\ ms$
20	$400\ \mu s$	$\approx 1.04\ s$
30	$900\ \mu s$	$\approx 18\ m$
40	$1.6\ ms$	$\approx 13\ d$
50	$2.5\ ms$	$\approx 36\ y$
60	$3.6\ ms$	$\approx 36535\ y$

Table 1: Complexity classes

1.4 Complexity classes

Two classes of algorithms are considered, according to their worst case order of complexity:

- **Polynomial:** $\mathcal{O}(n^d)$ for a constant $d > 0, d \in \mathbb{R}$
- **Exponential:** $\mathcal{O}(d^n)$ for a constant $d > 0, d \in \mathbb{R}$

Algorithms with a high order Polynomial complexity are not considered efficient.

A comparison of the two classes, assuming that $1\ \mu s$ is needed for each elementary operation, is shown in Table 1.

2 Graph and Network Optimization

Many **decision making problems** can be formulated in terms of graphs and networks, such as:

- **transportation** and **distribution** problems
- **network design** problems
- **location** problem
- timetable **scheduling**
- ...

2.1 Graphs

A **graph** is a pair $G = (N, E)$ with:

- N a set of **nodes** or **vertices**
- $E \subseteq N \times N$ a set of **edges** or arcs connecting them pairwise
 - A **directed** graph is a graph in which set E is composed by ordered pairs of distinct nodes
 - an edge connecting nodes i and j is represented by $\{i, j\}$
 - the flow is permitted only in the direction of the arrow (*from i to j*)
 - i is the head of the arc and j is the tail
 - A **undirected** graph is a graph in which set E is composed by unordered pairs of distinct nodes
 - an edge connecting nodes i and j is represented by (i, j)
 - since all edges are undirected, (i, j) is equivalent to (j, i)

Properties:

- Two **nodes** i and j are **adjacent** if they are connected by an edge (i, j) or $\{i, j\}$
- An **edge** e is **incident** in a node v if v is an **endpoint** of e
 - **undirected** graphs: the degree of a node is the number of incident edges
 - **directed** graphs: the in-degree (*out-degree*) of a node is the number of arcs that have it as successor (*predecessor*)
- A **path** from $i \in N$ to $j \in N$ is a sequence of edges in a undirected graph

$$p = \langle \{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{k-1}, v_k\} \rangle$$

connecting nodes v_1, \dots, v_k , with $\{v_i, v_{i+1}\} \in E$ for $i = 1, \dots, k-1$

- A **directed path** $i \in N$ to $j \in N$ is a sequence of arcs in a directed graph

$$p = \langle (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v_k) \rangle$$

connecting nodes, with $(v_i, v_{i+1}) \in E$ for $i = 1, \dots, k-1$

- A **cycle** is a path

$$\langle v_1, v_2, \dots, v_1 \rangle$$

where the first and last nodes are the same

- A **directed cycle** is path

$$\langle \{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{k-1}, v_k\}, \{v_k, v_1\} \rangle$$

where the first and the last nodes are the same

- Nodes u and v are **connected** if exists a path connecting them
- A graph (N, E) is **connecting** if u, v are connecting $\forall u, v \in N$
- A graph (N, E) is **strongly connected** if u, v are connected by a **directed** path $\forall u, v \in N$
- A graph is **bipartite** if there is a partition $N = N_1 \cup N_2$, $N_1 \cap N_2 = \emptyset$ such that $\forall (u, v) \in E, u \in N_1$ and $v \in N_2$
- A graph is **complete** if $E = \{\{v_i, v_j\} \mid v_i, v_j \in N \wedge i \leq j\}$

- Given a directed graph $G = (N, A)$ and $S \subseteq N$, the **outgoing cut** induced by S is the set of arcs:

$$\delta^+(S) = \{(u, v) \in A \mid u \in S \wedge v \in N \setminus S\}$$

the **incoming cut** induced by S is the set of arcs:

$$\delta^-(S) = \{(u, v) \in A \mid v \in S \wedge u \in N \setminus S\}$$

- A graph with n **nodes** has at most $m = \frac{n(n-1)}{2}$ **edges**
- A **directed** graph with n **nodes** has at most $m = n(n-1)$ **arcs**
 - a graph is **dense** if $m \approx n^2$
 - a graph is **sparse** if $m \ll n$
- The adjacency list $A(i)$ of a node i is the set of arcs emanating from that node

$$A(i) = \{j \in N \mid (i, j) \in A\}$$

Some examples are shown in Figure 2.

2.1.1 Graphs representation

Graphs are represented by:

- Adjacency **matrix** \mathbf{A} of size $n \times n$ if the graph is dense:

$$a_{ij} = \begin{cases} 1 & \text{if } (i, j) \in A \\ a_{ij} & \text{otherwise} \end{cases}$$

- Adjacency **list** \mathbf{A} of size n if the graph is sparse

The same representation can be used for both directed and undirected graphs; the adjacency matrix for an undirected graph is **symmetric**.

An example of a graph representation is shown in Figure 3.

2.1.2 Graph reachability problem

Definition 2.1 (Graph reachability problem). Given a directed graph $G = (N, A)$ and a node $s \in N$, find all nodes reachable from s .

Goal

- *Input*: graph $G = (N, A)$, described via successor lists, and a node $s \in N$
- *Output*: subset $M \subseteq N$ of nodes of G reachable from s

The goal is reached by an *efficient* algorithm to solve the problem, with the following properties:

- a **queue** Q of nodes not yet processed is kept by the algorithm
- the queue uses a **FIFO** policy
- the nodes exploration is performed in a **breadth-first** manner

Algorithm The algorithm pseudocode is shown in Code 1.

The algorithm stops when $\delta^+(M) = \emptyset$ (*when the outgoing cut of the set of nodes M is empty*); $\delta^-(M)$ is the set of arcs with head node in M and tail in $N \setminus M$.

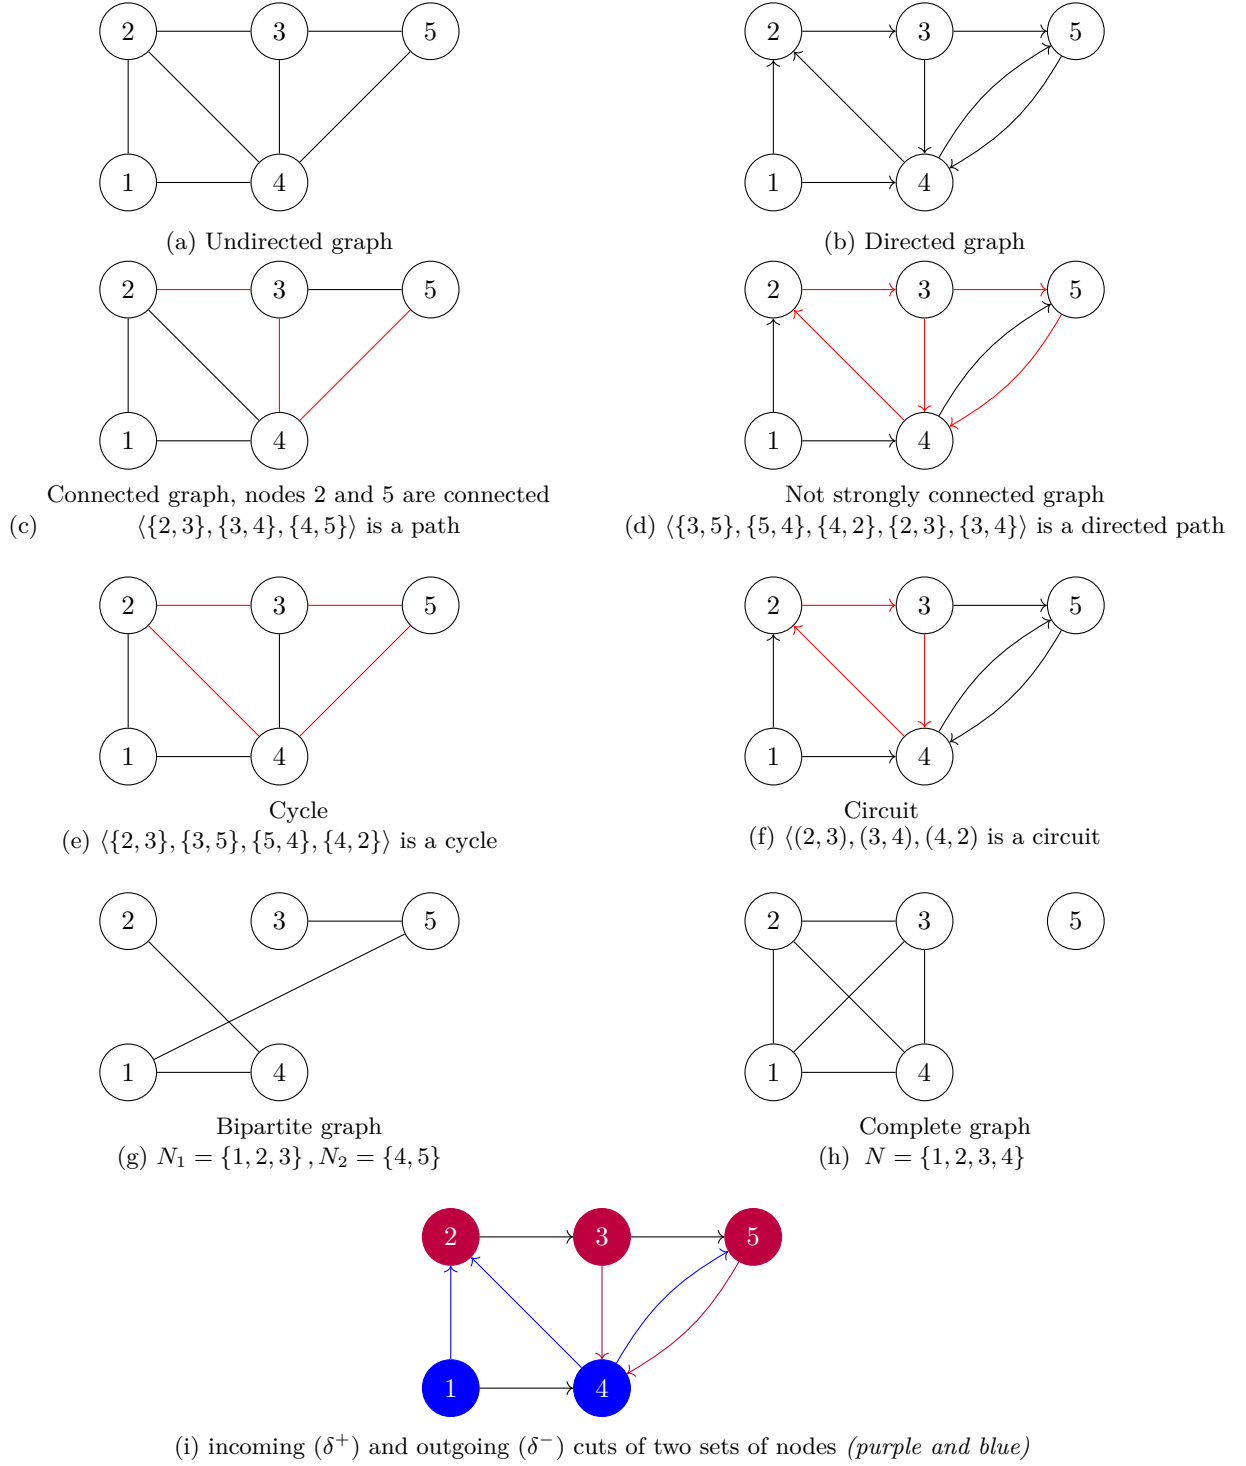


Figure 2: Examples of graphs

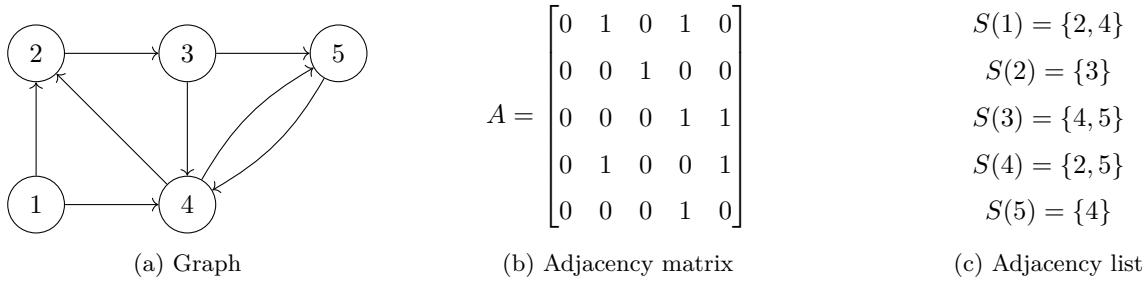


Figure 3: Graph representation

```

Q := {s}
M := {}
while Q is not empty do
  u := node ∈ Q
  Q := Q \ {u}
  M := M ∪ {u}
  for (u, v) ∈ δ+(u) do
    if v ∉ M and v ∉ Q then
      Q := Q ∪ {v}
    end
  end
end

```

Code 1: Graph reachability

2.1.2.1 Complexity analysis

At each iteration of the **while** loop:

1. A node u is **removed** from the queue Q and **added** to the set M
2. For all nodes v directly reachable from u and not already in M or Q , v is added to Q

Since each node u is inserted in Q at most once and each arch (u, v) is considered at most once, the overall complexity is:

$$\mathcal{O}(n + m) \quad n = |N|, m = |A|$$

For dense graphs, this value converges to $\mathcal{O}(n^2)$.

2.2 Subgraphs and Trees

Let $G = (N, E)$ be a graph. Then:

- $G' = (N', E')$ is a **subgraph** of G if $N' \subseteq N$ and $E' \subseteq E$
- A **tree** $G_T = (N', T)$ of G is a connected, acyclic, subgraph of G
- $G_T = (N', T)$ is a **spanning tree** of G if it contains all the nodes ($N' = N$)
- The **leaves** of a tree are the nodes with degree 1

A representation of these concepts is shown in Figure 4.

2.3 Properties of trees

Property 2.1 (Number of edges). Every tree with n nodes has $n - 1$ edges.

Proof.




Figure 4: Subgraphs and trees

- **Base case:** the claim holds for $n = 1$ (a tree with a single node has no edges)
- **Inductive steps:** show that the claim is valid for for a tree with $n + 1$ nodes
 - let T_1 be a tree with $n + 1$ and recall with any tree with $n \geq 2$ nodes has at least 2 leaves
 - by deleting one of the leaves and its incident edge, a tree T_2 with n nodes is obtained
 - by induction hypothesis, T_2 has $n - 1$ edges; therefore, T_1 has $n - 1 + 1 = n$ edges

□

Property 2.2. Any pair of nodes in a tree is connected via a unique path. Otherwise, the tree would contain a cycle.

Property 2.3. By adding a new edge to a tree, a new unique cycle is created. This cycle consists of the path created in Property 2.2 and the new edge.

Property 2.4. Let $G_T = (N, T)$ be a spanning tree of $G = (N, E)$. Consider an edge $e \notin T$ and the unique cycle C of $T \cup \{e\}$. For each edge $f \in C \setminus \{e\}$, the subgraph $T \cup \{e\} \setminus \{f\}$ is a spanning tree of G .




Figure 5: Exchange property

Property 2.5. Let F be a partial tree (spanning nodes in $S \subseteq N$) contained in a optimal spanning tree of $G = (N, E)$. Consider $e = \{u, v\} \in \delta(S)$ of minimum cost, then there exists a minimum cost spanning tree of G containing e .

Proof. By contradiction, assume $T^* \subseteq E$ is a minimum cost spanning tree with $F \subseteq T^*$ and $e \notin T^*$. Adding an edge e to T^* creates the cycle C . Let $f \in \delta(S) \cap C$:

- If $c_e = c_f$, then $T^* \cup \{e\} \setminus \{f\}$ is a minimum cost spanning tree of G as it has the same cost as T^*
- If $c_e < c_f$, then $c(T^* \cup \{e\} \setminus \{f\}) < c(T^*)$, hence T^* is not optimal

□

2.4 Optimal cost spanning tree

Spanning trees have a number of applications, including:

- **network** design
- **IP network** protocols
- **compact memory** storage

Model: an undirected graph $G = (N, E)$, $n = |N|$, $m = |E|$ and a cost function $c : E \rightarrow \mathbb{R}$, that assigns a cost to each edge, with $e = \{u, v\} \in E$.

Required properties

1. Each **pair of nodes** must be in a **path**
 \Rightarrow the output must be a **connected subgraph** containing all the nodes N of G
2. The **subgraph** must have **no cycles**
 \Rightarrow the output must be a **tree**

Definition 2.2 (Problem definition). Given an undirected graph $G = (N, E)$ and a cost function $c : E \rightarrow \mathbb{R}$, find a spanning tree $G_T(N, T)$ of G of minimum, total cost.

The objective is finding:

$$\min_{T \in X} \sum_{e \in T} c_e \quad X = \text{set of all spanning trees of } G$$

Theorem 2.1. A complete graph with n nodes ($n \geq 1$) has n^{n-2} spanning trees.

Property 2.6. Every spanning tree of a connected n -node graph has $n - 1$ edges.

2.4.1 Prim's algorithm

Idea: iteratively build a spanning tree.

Method

1. Start from initial tree (S, T) with $S = \{u\}$, $S \subseteq N$ and $T = \emptyset$
2. At each step, add to the current partial tree (S, T) an edge of minimum cost among those which connect a node in S to a node in $N \setminus S$

Goal

- \rightarrow *Input:* connected graph $G = (N, E)$ with edge costs.
- \rightarrow *Output:* subset $T \subseteq E$ of edges of G such that $G_T = (N, T)$ is a minimum cost spanning tree of G .

Complexity if all edges are scanned at each iteration, the complexity order is $\mathcal{O}(nm)$

Algorithm the pseudocode of the algorithm is shown in Code 2. Prim's algorithm is **greedy**: at each step a minimum cost edge is selected among those in the cut $\delta(S)$ induced by the current set of nodes S .

2.4.1.1 Correctness of Prim's algorithm

Proposition 2.1. Prim's algorithm is exact.

The exactness does not depend on the choice of the first node nor on the selected edge of minimum cost $\delta(S)$. Each selected edge is part of the optimal solution as it belongs to a minimum spanning tree.

The optimality condition allows to verify whether a spanning tree T is optimal or not; it suffices to check that each $e \in E \setminus T$ is not a cost decreasing edge.

```

S := {u}
T := {}
while |T| < n - 1 do
  {u, v} := edge ∈ δ(S) with minimum cost // u ∈ S, v ∈ N \ S
  S := S ∪ {v}
  T := T ∪ {{u, v}}
end

```

Code 2: Prim's algorithm

2.4.1.2 Implementation in quadratic time

The Prim's algorithm can be implemented in quadratic time, i.e. $\mathcal{O}(n^2)$.

Data structure

- k number of edges selected so far
- Subset $S \subseteq N$ of nodes incident to the selected edges
- Subset $T \subseteq E$ of selected edges
- $C_j = \begin{cases} \min\{c_{ij} \mid i \in S\} & j \notin S \\ +\infty & \text{otherwise} \end{cases}$
- $closest_j = \begin{cases} \arg \min\{c_{ij} \mid i \in S\} & j \notin S \\ \text{predecessor of } j \text{ in the minimum spanning tree} & j \in S \end{cases}$

An example of a step is shown in Figure 6.




Figure 6: Data structure
nodes 1, 2, 3 ∈ S, node j ∉ S
 $closest_j = 3$ $c_{closest_j j} = 1$

The spanning tree is built by selecting the node j with minimum cost C_j and adding the edge $\{j, closest_j\}$ to the spanning tree.

The code for this algorithm is shown in Code 3.

The complexity of this algorithm is $\mathcal{O}(n^2)$. For sparse graphs, where $m \ll \frac{n(n-1)}{2}$, a more efficient implementation ($\mathcal{O}(m \log(n))$) (using priority queues) is possible.

2.4.2 Optimality condition

Given a spanning tree T , an edge $e \notin T$ is **cost decreasing** if when added to T , it creates a cycle C with $C \subseteq T \cup \{e\}$ and $\exists f \in C \setminus \{e\}$ such that $c_e < c_f$.

Theorem 2.2. A tree T is of minimum total cost if and only if no cost decreasing edge exists.

```

T := {}
S := {u}
// initialization
for j  $\notin$  N \ S do
    if {u, j}  $\in$  E then
        C_j := c_{u, j}
    else
        C_j := +\infty
    end
    closest_j := u
end
for k := 1 to n - 1 do
    min := + $\infty$  // selection of min cost edge
    for j := 1, ..., n do
        if j  $\notin$  S and C_j < min then
            min := C_j
            v := j
        end
    end
    S := S  $\cup$  {v} // extend S
    T := T  $\cup$  {{v, closest_v}} // extend T
    for j := 1 to n do
        if j  $\notin$  S and c_vj < C_j then
            C_j := c_vj
            closest_j := v
        end
    end
end
end

```

Code 3: Prim's algorithm in quadratic time

Proof.

- \Rightarrow If a **cost decreasing edge exists**, then T is **not of minimum total cost**
- \Leftarrow If **no cost decreasing edge exists**, then T is **of minimum total cost**
 - let T^* be a minimum cost spanning tree of graph G , found via by Prim's algorithm
 - it can be verified that T^* can be iteratively (*changing one edge at a time*) transformed into T without changing the total cost
 - thus, T is also optimal

□

2.5 Optimal paths

Optimal (*shortest, longest, ...*) paths have a wide range of applications, including:

- **Google Maps, GPS navigators**
- Planning and management of **transportation, electrical, and telecommunication networks**
- **Problem** planning

Model

Given a directed graph $G = (N, A)$ with a cost $c_{ij} \in \mathbb{R}$ associated to each arc $(i, j) \in A$, and two nodes s and t , determine a minimum cost (*shortest*) path from s to t .

- Each **value** $c_{i,j}$ represents the **cost** (*or length, travel time, ...*) of arc $(i, j) \in A$
- Node s is the **origin** (*or source*), node t is the **destination** (*or sink*)

Property 2.7. A path $\langle (i_1, i_2), (i_2, i_3), \dots, (i_{k-1}, i_k) \rangle$ is **simple** if no node is visited more than once

Property 2.8. If $c_{ij} \geq 0$ for all $(i, j) \in A$, there is at least one shortest path that is simple.

2.5.1 Dijkstra's algorithm

Idea: consider the nodes in increasing order of length (*cost*) of the shortest path from s to any one of the other nodes.

Method

- To each **node** $j \in N$, a **label** L_j is associated
 - \Rightarrow at the end of the algorithm, this label will be the cost of the minimum cost path from s to j
- Another label $predecessor_j$ is associated to each node $j \in N$
 - \Rightarrow at the end of the algorithm, this label will be the node that precedes j on the minimum cost path from s to j
- Make a **greedy** choice with respect to the paths from s to j
- A set of **shortest paths** from s to any node $j \notin s$ can be retrieved backwards from t to s iterating over the predecessors

Goal

- \rightarrow *Input:* graph $G = (N, A)$, cost $c_{ij} \geq 0 \forall i, j$, origin $s \in N$
- \rightarrow *Output:* shortest path from s to all other nodes in G

Data structure

- $S \subseteq N$: subset of nodes whose labels are permanent
- $X \subseteq N$: subset of nodes with temporary labels

```

S := {}
X := {s}
for u ∈ N do
    L_u := ∞
end
L_s := 0
while |S| < |N| do
    u := argmin{L_i | i ∈ X}
    X := X \ {u}
    S := S ∪ {u}
    for (u, v) ∈ δ+(u) do
        if L_v > L_u + c_uv then
            L_v := L_u + c_uv
            predecessor_v := u
            X := X ∪ {v}
        end
    end
end

```

Code 4: Dijkstra's algorithm

- $L_j = \begin{cases} \text{cost of a shortest path from } s \text{ to } j & j \in S \\ \min\{L_i + c_{ij} \mid (i, j) \in \delta^+(S)\} & j \notin S \end{cases}$
 - given a directed graph G and the current subset of nodes $S \subset N$, consider the outgoing cut $\delta^+(S)$ and select $(u, v) \in \delta^+(S)$ such that: $L_u + c_{uv} = \min\{L_i + c_{ij} \mid (i, j) \in \delta^+(S)\}$
 - thus: $L_u + c_{uv} \leq L_i + c_{ij}, \forall (i, j) \in \delta^+(S)$
- $\text{predecessor}_j = \begin{cases} \text{predecessor of } j \text{ in the shortest path from } s \text{ to } j & j \in S \\ u \text{ such that } L_u + c_{uj} = \min\{L_i + c_{ij} \mid i \in S\} & j \notin S \end{cases}$

Complexity: the complexity of the algorithm depends on the how the arc (u, v) is selected among those of the current cut $\delta^+(u)$.

- If all m arcs are scanned, the overall complexity would be $\mathcal{O}(nm)$, hence $\mathcal{O}(n^3)$
- If all labels L_j are determined by appropriate updates (as in Prim's algorithm), only a single arc of $\delta^+(j)$ is scanned, hence the complexity is $\mathcal{O}(n^2)$

Notes:

- A set of shortest paths from s to all the nodes $j \in N$ can be retrieved backwards from t to s iterating over the predecessors
- The union of a set of shortest paths from node s to all the other nodes of G is an arborescence rooted at s
- Dijkstra's algorithm does not work when there are arcs with negative cost: if G contains a circuit of negative cost, the shortest path problem may not be well defined

The code for this algorithm is shown in Code 4.

2.5.1.1 Correctness of Dijkstra's algorithm

Proposition 2.2. Dijkstra's algorithm is correct.

Proof.

1. At the k -th step:

- $S = \{s, i_1, \dots, i_{k-1}\}$
- $\begin{cases} \text{cost of a minimum cost path from } s \text{ to } j & j \in S \\ \text{cost of a minimum cost path with all intermediate nodes in } S & j \notin S \end{cases}$

2. By induction on the number k of steps:

- base case: for $k = 1$ the statement holds, since

$$S = \{s\}, \quad L_s = 0, \quad L_j = +\infty, \quad \forall j \notin S$$

- inductive step: assume that the statement holds for $k + 1$

- let $u \notin S$ be the node that is inserted in S and ϕ the path from s to u such that:

$$L_v + c_{vu} \leq L_i + c_{iu}, \quad \forall (i, v) \in \delta^+(S)$$

- every path π from s to u has $c(\pi) \geq c(\phi)$, as there exists $i \in S$ and $j \notin S$ such that:

$$\pi = \pi_1 \cup \{(i, j)\} \cup \pi_2$$

where (i, j) is the first arc in $\pi \cap \delta^+(S)$

- it holds that

$$c(\pi) = c(\pi_1) + c_{ij} + c(\pi_2) \geq L_i + c_{ij}$$

because $c_{ij} \geq 0 \Rightarrow c(\pi_2) \geq 0$ and by the choice of (v, u) , $c(\pi_1) \geq L_i$

- finally, by induction assumption:

$$L_i + c_{ij} \geq L_v + c_{vu} = c(\phi)$$

- a visualization of this step of the proof is shown in Figure 7

□






Figure 7: Proof of the induction step; nodes s, v, i are in cut S

2.5.1.2 Example of Dijkstra's algorithm


An example of Dijkstra's algorithm is shown in Figure 8.




(a) Sample graph, with the cost of each arc




(b) step 1 of Dijkstra's algorithm




(c) step 2 of Dijkstra's algorithm




(d) step 3 of Dijkstra's algorithm



(e) step 4 of Dijkstra's algorithm



(f) step 5 of Dijkstra's algorithm



(g) step 6 of Dijkstra's algorithm

Figure 8: Example of Dijkstra's algorithm

2.5.2 Floyd-Warshall's algorithm

Goal

- *Input*: a directed graph $G = (N, A)$ with an $n \times n$ cost matrix $C = [c_{ij}]$
- *Output*: for each pair of nodes $i, j \in N$, the cost c_{ij} of the shortest path from i to j

Data structure

- Two $n \times n$ matrices D, P whose elements correspond, at the end of the algorithm, to:
 - d_{ij} the cost of the shortest path from i to j
 - p_{ij} the predecessor of j on the shortest path from i to j

Method

1. Initialization of D and P :

$$p_{ij} = i \quad \forall i$$

$$d_{ij} = \begin{cases} 0 & i = j \\ c_{ij} & i \neq j \wedge (i, j) \in A \\ +\infty & \text{otherwise} \end{cases}$$

2. Triangular operation: for each pair of nodes i, j , where $i \neq u, j \neq u$, check whether the path from i to j is shorter by going through u (i.e. $d_{iu} + d_{uj} < d_{ij}$)

Complexity

- Since in the worst case the triangular operation is executed for all nodes u and for each pair of nodes i, j , the complexity is $\mathcal{O}(n^3)$

The code for this algorithm is shown in Code 5.

2.5.2.1 Correctness of Floyd-Warshall's algorithm

Proposition 2.3. Floyd-Warshall's algorithm is correct.

Proof. assume that the nodes of G are numbered from 1 to n . Verify that, if the node index order is followed, after the u -th cycle the value d_{ij} (for any i, j) corresponds to the cost of a shortest path from i to j with at most u intermediate nodes ($\{1, \dots, u\}$) □

2.6 Optimal paths in directed, acyclic graphs

A directed graph $G = (N, A)$ is **acyclic** if it does not contain any circuit. A directed acyclic graph G is then referred to as a **DAG**.

Property of **DAGs**: the nodes of any directed acyclic graph G can be ordered topologically, i.e. indexed so that for each arc $(i, j) \in A$ the index of i is less than the index of j ($i \leq j$). The topological order can be exploited by dynamic programming algorithms to compute efficiently the shortest paths in a **DAG**.

Problem: given a **DAG** $G = (N, A)$ with a cost $c_{ij} \in \mathbb{R}$ and nodes s, t , determine the shortest (or longest) path from s to t .


```

for j := 1 to n do
  p_id := i


  if i = j then
    d_ij := 0
  else if (i, j) in A then
    d_ij := c_ij
  else
    d_ij := +∞
  end
end
end
for u ∈ N do
  for i ∈ N \{ u } do
    for j ∈ N \{ u }
      if d_iu + d_uj < d_ij then
        p_ij := p_uj
        d_ij := d_iu + d_uj
      end
    end
  for i ∈ N do
    if d_ij < 0 then
      error "negative cycle"
    end
  end
end

```

Code 5: Floyd-Warshall's algorithm



(a) Topological ordering



(b) Node v with $\delta^-(v) = \emptyset$, as in step (1) of the algorithm

Figure 9: Topological ordering method

2.6.1 Topological ordering method

The method requires $G = (N, A)$ to be a *DAG* represented via the list of predecessors $\delta^-(v)$ and the list of successors $\delta^+(v)$ of each node $v \in N$. Then, it works as follows:

1. Assign the smallest positive integer not yet assigned to a node $v \in N$ with $\delta^-(v) = \emptyset$
 \rightarrow such node always exists because G does not contain circuits
2. Delete the node v with all its incident arcs
3. Go to step (1) until all nodes have been assigned a number

This algorithm has complexity $\mathcal{O}(|A|)$, because each node is assigned a number only once. Furthermore, all arcs incident to a node are deleted only once.

2.6.2 Dynamic programming for shortest path in *DAGs*

Any shortest path from 1 to t , called π_t , with at least 2 arcs can be subdivided into two parts:

- π_i , the shortest subpath from s to i
- (i, t) , the remaining part

This decomposition is called the optimality principle of shortest paths in *DAGs*. An illustration of this decomposition is shown in Figure 10.



Figure 10: Shortest path from 1 to t

The strategy to find the shortest path is:

1. For each node $i = 1, \dots, t$ let L_i be the cost of a shortest path from 1 to i
 $\rightarrow L_t = \min_{(i,t) \in \delta^-(t)} \{L_i + c_{it}\}$
 \rightarrow the minimum is taken over all possible predecessors i of t
2. If G is topologically ordered *DAG*, then the only possible predecessors of t in a shortest path π_t from 1 to t are those with index $i < t$
 $\rightarrow L_t = \min_{i < t} \{L_i + c_{it}\}$
 \rightarrow in a graph with circuits, any node i can be a predecessor of t if $i \neq t$

For *DAGs* whose nodes are topologically ordered L_{t-1}, \dots, L_1 satisfy the same type of recursive relations:

$$L_{t-1} = \min_{i < t-1} \{L_i + c_{i,t-1}\}; \dots; L_2 = \min_{i=1} \{L_i + c_{i2}\} = L_1 + c_{12}; L_1 = 0$$

which can be solved in reversed order

$$L_1 = 0; L_2 = L_1 + c_{12}; \dots; L_t = \min_{i < t-1} \{L_i + c_{it}\}$$

Algorithm: finally, the algorithm is shown in pseudocode in Code 6.

Complexity of the algorithm is $\mathcal{O}(|A|)$:

- Topological ordering of the nodes: $\mathcal{O}(m)$ with $m = |A|$ (number of arcs)
- Each node/arc is processed only once: $\mathcal{O}(n + m)$

In order to find the longest path, the algorithm can be adapted as follows:

$$L_t = \max_{i < t} \{L_i + c_{it}\}$$

```

sort the nodes of G topologically
L_1 := 0
for j := 2 to n do
  L_j := min{L_i + c_{ij} | (i, j) ∈ δ-(j) ∧ i < j}
  pred_j := v such that (v, j) = argmin{L_i + c_{ij} | (i, j) ∈ δ-(j) ∧ i < j}
end

```

Code 6: Shortest path in DAG

2.6.2.1 Optimality of the algorithm

The Dynamic Programming algorithm for finding shortest or longest paths in DAGs is exact. This is due to the optimality principle, already explored in the previous section.

2.7 Project planning

A project consists of a set of m activities with their (estimated) duration: activity A_i has duration $d_i \geq 0, i = 1, \dots, m$. Some pair of activities allow a precedence constraint: $A_i \propto A_j$ indicated that A_i must be performed before A_j .

Model: a project can be represented by a directed graph $G = (N, A)$ where:

- each arc corresponds to an activity
- the arc length represent the duration of the corresponding activity

In order to account for precedence constraints, the arcs must be positioned such that for activities $A_i \propto A_j$ there exists a directed path where the arc associated to A_i precedes the arc associated to A_j . Such notation is shown in Figure 11.




Figure 11: Precedence relation in project planning

Therefore, a node v marks an event corresponding to the end of all the activities $(i, v) \in \delta^-(v)$ and the (possible) start of all the activities $(v, j) \in \delta^+(v)$.

Property 2.9. The directed graph G representing a project is acyclic (is a DAG).

Proof. by contradiction, if $A_{i1} \propto A_{i2} \propto \dots \propto A_{jk} \propto A_{kj}$ there would be a logical inconsistency. □

2.7.1 Optimal paths

A graph G can be simplified by contracting some arcs, but it's important to not introduce unwanted precedence constraints. Artificial nodes or artificial arcs are introduced so that graph G :

- Contains a unique initial node s corresponding to the event “beginning of the project”
- Contains a unique final node t corresponding to the event “end of the project”
- Does not contain multiple arcs with the same origin and destination

Problem: given a project (set of activities with duration and precedence constraints), schedule the activities in order to minimize the overall project duration (the time needed to complete all the activities).

```

sort the nodes topologically
T_min_i := 0
for j = 2 to n do
    T_min_j := max{T_min_i + d_ij | (i, j) ∈ δ-(j)}
end for
T_max_n := T_min_n // minimum project duration
for i = n-1 to 1 do
    T_max_i := min{T_max_j - d_ij | (i, j) ∈ δ+(i)}
end for

```

Code 7: Critical path method

Property 2.10. The minimum overall project duration is the length of a longest path from s to t in the graph G .

Proof. since any $s - t$ path represents a sequence of activities that must be executed in the specified order, its length provides a lower bound on the minimum overall project duration. \square

2.7.1.1 Critical path method - CPM

The critical path method (CPM) determines:

- A **schedule** (a plan for executing the activities specifying the order and the assigned time) that minimizes the overall project duration
- The **slack** of each activity (the amount of time by which its execution can be delayed without affecting the overall minimum project duration)

Initialization construct the graph G representing the project.

Method

1. Find a topological order of the nodes
2. Consider the nodes by increasing indices and for each $h \in N$ find the earliest time T_{min_h} at which the event associated to node h can occur
 $\rightarrow T_{min_h}$ corresponds to the minimum project duration
3. Consider the nodes by decreasing indices and for each $h \in N$ find the latest time T_{max_h} at which the event associated to node h can occur without delaying the project completion date beyond T_{min_n}
4. For each activity $(i, j) \in A$ find the slack
 \rightarrow the slack is calculated as $\sigma_{ij} = T_{max_j} - T_{min_i} - d_{ij}$

Goal

- \rightarrow *Input:* graph $G = (N, A)$ with $n = |N|$ and the duration d_{ij} associated to each $(i, j) \in A$
- \rightarrow *Output:* $(T_{min_i}, T_{max_i}), i = 1, \dots, n$

Algorithm: finally, the algorithm is shown in pseudocode in Code 7.

Complexity: the overall complexity is $\mathcal{O}(n + m) \approx \mathcal{O}(m)$, due to the sum of

- complexity of the topological sort - $\mathcal{O}(n + m)$
- complexity of the first loop - $\mathcal{O}(n + m)$
- complexity of the second loop - $\mathcal{O}(n + m)$

2.7.1.2 Critical paths

An activity (i, j) with zero slack $\sigma_{ij} = T_{max_j} - T_{min_i} - d_{ij} = 0$ is called **critical**.

A critical path is a path in a $s - t$ composed uniquely by critical activities. At least one always exists.

2.7.1.3 Gantt charts

A Gantt chart is a graphical representation of a project schedule. It was introduced in 1896 by Henry Gantt, an American mechanical engineer and management consultant.

There are two types of Gantt charts:

- Gantt chart at **earliest** - each activity (i, j) starts at T_{min_i} and ends at $T_{min_i} + d_{ij}$
- Gantt chart at **latest** - each activity (i, j) starts at T_{max_i} and ends at $T_{max_i} + d_{ij}$

3 Linear Programming

3.1 Optimization problems

Optimization problems are problems that require to find the best solution among a set of possible solutions.

Definition 3.1. An instance of an optimization problem is a pair (F, c) where:

- F is the domain of feasible point
- c is the cost function, a mapping $c : F \rightarrow \mathbb{R}$
- the problem is finding an $f \in F$ such that $c(f) \leq c(y) \forall y \in F$

Such point is called a globally optimal (*or just optimal*) solution to the given instance.

Definition 3.2. An **optimization problem** is a set of I instances of a given optimization problem.

Definition 3.3. A **linear programming** (*or LP*) problem is an optimization problem such as

$$\begin{aligned} \min f(x) \\ \text{s.t. } x \in X \subseteq \mathbb{R}^n \rightarrow \mathbb{R} \end{aligned} \quad (3.1)$$

where:

- the **objective function** $f : X \rightarrow \mathbb{R}$ is **linear**
- the **feasible region** $X = \{x \in \mathbb{R}^n \mid g_i(x) r_i 0 \wedge i \in \{1, \dots, m\}\}$ with $r_i \in \{=, \geq, \leq\}$ and $g_i : \mathbb{R}^n \rightarrow \mathbb{R}$ are **linear** functions for $i \in \{1, \dots, m\}$
- $x^* \in \mathbb{R}^n$ is an **optimal solution** of the LP 3.1 if $f(x^*) \leq f(x) \forall x \in X$

A wide variety of decision making problems can be formulated or approximated as *LP*, as they often involve the optimal allocation of a given set of limited resources to different activities.

- **General form** of a linear programming problem:

$$\begin{aligned} \min z &= c_1 x^1 + \dots + c_n x_n \\ \text{s.t. } a_{11} x^1 + \dots + a_{1n} x_n &(\leq, =, \geq) b_1 \\ &\vdots \\ a_{m1} x^1 + \dots + a_{mn} x_n &(\leq, =, \geq) b_m \\ x^1, \dots, x_n &\geq 0 \end{aligned} \quad (3.2)$$

- **Matrix notation** of a linear programming problem:

$$\begin{aligned} \min z &= \begin{bmatrix} c_1 & c_2 & \dots & c_n \end{bmatrix} \begin{bmatrix} x^1 \\ x^2 \\ \vdots \\ x_n \end{bmatrix} \\ \text{s.t. } \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} x^1 \\ x^2 \\ \vdots \\ x_n \end{bmatrix} &\{ \leq, =, \geq \} \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_m \end{bmatrix} \\ \begin{bmatrix} x^1 \\ \vdots \\ x_n \end{bmatrix} &\geq 0 \end{aligned} \quad (3.3)$$

3.2 Assumptions of LP models

The LP model is based on the following **assumptions**:

- **Linearity** (*proportionality and additivity*) of the objective function and constraints
 - **proportionality**: *contribution of each variable = constant × variable*. It does not account for economies of scale
 - **additivity**: *total contribution = \sum_i contribution of each variable i* . It does not account for competing activities (their sum is not necessarily the total contribution)
- **Divisibility** of the variables, as they can assume fractional (rational) values
- **Parameters** are assumed to be **constants** which can be estimated with a sufficient degree of accuracy
 - more complex mathematical programs are needed to account for uncertainty in the parameter values

LP **sensitivity analysis** allows to evaluate how “*sensitive*” an optimal solution is with respect to small changes in the parameters of the model.

3.3 Equivalent Forms

The **General Form** (or *canonical form*) (Equation 3.4) of a LP can be expressed in the equivalent **Standard Form** (Equation 3.5).

$$\begin{aligned}
 \min(\max) \quad & z = c^T x \\
 \text{s.t.} \quad & A_1 x \geq b_1 && \text{inequality constraints} \\
 & A_2 x \leq b_2 && \text{inequality constraints} \\
 & A_3 x = b_3 && \text{equality constraints} \\
 & x_j \geq 0, \quad j \in J \subseteq \{1, \dots, n\} && \text{non-negativity constraints} \\
 & x_j \text{ free}, \quad j \in \{1, \dots, n\} \setminus J && \text{free variables}
 \end{aligned} \tag{3.4}$$

$$\begin{aligned}
 \min \quad & z = c^T x \\
 \text{s.t.} \quad & Ax = b && \text{inequality constraints} \\
 & x \geq 0 && \text{non negative variables}
 \end{aligned} \tag{3.5}$$

The two forms are equivalent, as simple transportation rules allow to pass from one form to the other; the transformation may involve adding and or deleting variables and constraints, as the next Section shows.

3.3.1 Transformation rules

- $\max c^T x \Rightarrow \min -c^T x$
- $a^T x \leq b \Rightarrow \begin{cases} a^T x + s = b \\ s \geq 0 \end{cases}$ s is a slack variable
- $a^T x \geq b \Rightarrow \begin{cases} a^T x - s = b \\ s \geq 0 \end{cases}$ s is a surplus variable
- x_j unrestricted in sign $\Rightarrow \begin{cases} x_j = x_j^+ - x_j^- \\ x_j^+ \geq 0 \\ x_j^- \geq 0 \end{cases}$

→ after the substitution, x_j is deleted from the problem

- $a^T x \leq q \Leftrightarrow -a^T x \geq -q$
- $a^T x \geq q \Leftrightarrow -a^T x \leq -q$
- $a^T x = b \Leftrightarrow \begin{cases} a^T x \leq b \\ a^T x \geq b \end{cases} \Leftrightarrow \begin{cases} a^T x \leq b \\ -a^T x \geq -b \end{cases}$

3.4 Graphical solutions

A **level curve** of value z of a function f is the set of points in \mathbb{R}^n where f is constant and takes value z .

Consider a LP with inequality constraints (*as it's easier to visualize*).

- A **hyperplane** is the set of points that satisfy the constraint $H = \{x \in \mathbb{R}^n \mid a^T x = b\}$
- An **affine half space** is the set of points that satisfies the constraint $H = \{x \in \mathbb{R}^n \mid a^T x \leq b\}$
 - each inequality constraint $a^T x \leq b$ defines an affine half space in the variable space
- The **feasible region** X of any LP is a polyhedron P defined by the intersection of a finite number of affine half spaces
 - P can be empty or unbounded
- A subset $S \subseteq \mathbb{R}^n$ is **convex** if for each pair of points $y^1, y^2 \in S$ the line segment between y^1 and y^2 is contained in S
 - given two points $y^1, y^2 \in S$, a **convex combination** of them is any point of the form

$$z = \lambda y^1 + (1 - \lambda) y^2 \quad 0 \leq \lambda \leq 1$$

- a convex combination with $\lambda \neq 0, 1$ is called **strict**
- The segment defined by $y^1, y^2 \in S$ defined by all the convex combinations of y^1 and y^2 is called a **convex hull**
 - $[y^1, y^2] = \{x \in \mathbb{R}^n \mid x = \alpha y^1 + (1 - \alpha) y^2 \wedge \alpha \in [0, 1]\}$
- A **polyhedron** P is a convex set of \mathbb{R}^n
 - any half space is convex
 - the intersection of a finite number of convex sets is also a convex set
- A **vertex of polyhedron** P is a point of P that cannot be expressed as a convex combination of other points of P
 - mathematically, x is a vertex P iff

$$x = \alpha y^1 + (1 - \alpha) y^2, \alpha \in [0, 1] \quad y^1, y^2 \in P \Rightarrow x = y^1 \vee x = y^2$$

- a non empty polyhedron $P = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$ has a finite number ($n > 1$) of vertices
- A **polytope** is a bounded polyhedron
- Given a polyhedron P , a vector $d \in \mathbb{R}^n, d \neq \bar{0}$ is an **unbounded feasible direction** of P if, for every point $x^0 \in P$, the ray $\{x \in \mathbb{R}^n \mid x = x^0 + \lambda d, \lambda \geq 0\}$ is contained in P

3.4.1 Weyl-Minkowski Theorem

The *Weyl-Minkowski* Theorem on the representation of polyhedra states that:


Theorem 3.1. *Every point x of a polyhedron P can be expressed as a convex combination of its vertices x^1, \dots, x^k plus (if needed) an unbounded feasible direction d of P :*

$$x = \alpha_1 x^1 + \dots + \alpha_k x^k + d$$


where the multipliers $\alpha_i \geq 0$ satisfy the constraint $\sum_{i=1}^k \alpha_i = 1$.

The unbounded feasible direction is needed if the polyhedron is unbounded; Figure 13 represent the cases of a bounded and unbounded polyhedron.


Consequences: every point x of polytope P can be expressed as a convex combination of its vertices. Then the *Weyl-Minkowski theorem* can be used to describe any point. An example of this is shown in Figure 14.




(a) Hyperplane



(b) Affine half space




(c) Convex polyhedron




(d) Not convex polyhedron

Figure 12: Illustrations of LP geometry definitions



(a) Unbounded polyhedron



(b) Bounded polyhedron

Figure 13: Illustration of the *Weyl-Minkowski* theorem




Figure 14: Example of polytope. $x = \alpha_1 x^1 + \alpha_2 x^2 + \alpha_3 x^3$ with $\alpha_1 + \alpha_2 + \alpha_3 = 1, \alpha_i \geq 0, d = 0$

3.4.2 Geometry of LP

Geometrically:

- An **interior point** $x \in P$ cannot be an optimal solution of the problem
 - it always exists an improving direction
 - consider Figure 15a where c represents the direction of fastest increase in z (*constant gradient*)
- In an optimal vertex, all **feasible direction are worsening directions**
 - consider Figure 15b where c represents the direction of fastest increase in z (*constant gradient*)
- The *Weyl-Minkowski* theorem implies that, although the variables can assume fractional values, any LP can be seen as a **combinatorial problem**
 - “only” the vertices of the polyhedron have to be considered in order to find the feasible solutions
 - the graphical method is only applicable for $n \leq 3$
 - the number of vertices often grows exponentially with respect to the number of variables

Furthermore, let P be a convex polytope of dimension d and let HS be a half space of P defined by hyperplane H . If the intersection $f = P \cap HS$ is a subset of H then:

- f is a **face** of P
- H is the **support hyperplane** of f

The face is then defined according to its dimension:

- a **vertex** is a face of dimension 0 (*a point*)
- a **facet** is a face of dimension $d - 1$ (*a hyperplane*)
- an **edge** is a face of dimension 1 (*a line*)

3.4.3 Four types of LP

There are four types of LP, depending on the number of solutions; all are illustrated in Figure 16. Since the objective of the problem is to minimize $f(x)$ (as it's in form $\min c^T x$), better solutions are found by moving along the direction $-c$ (*the opposite of the gradient $\nabla f(x)$*).

1. **A unique** optimal solution, Figure 16a
2. **Multiple** (*infinitely many*) optimal solutions, Figure 16b
3. **Unbounded LP**, Figure 16c
 - this type of problem has unbounded polyhedron and unlimited objective function values
4. **Infeasible LP**, Figure 16d
 - this type of problem has an empty polyhedron and no feasible solution

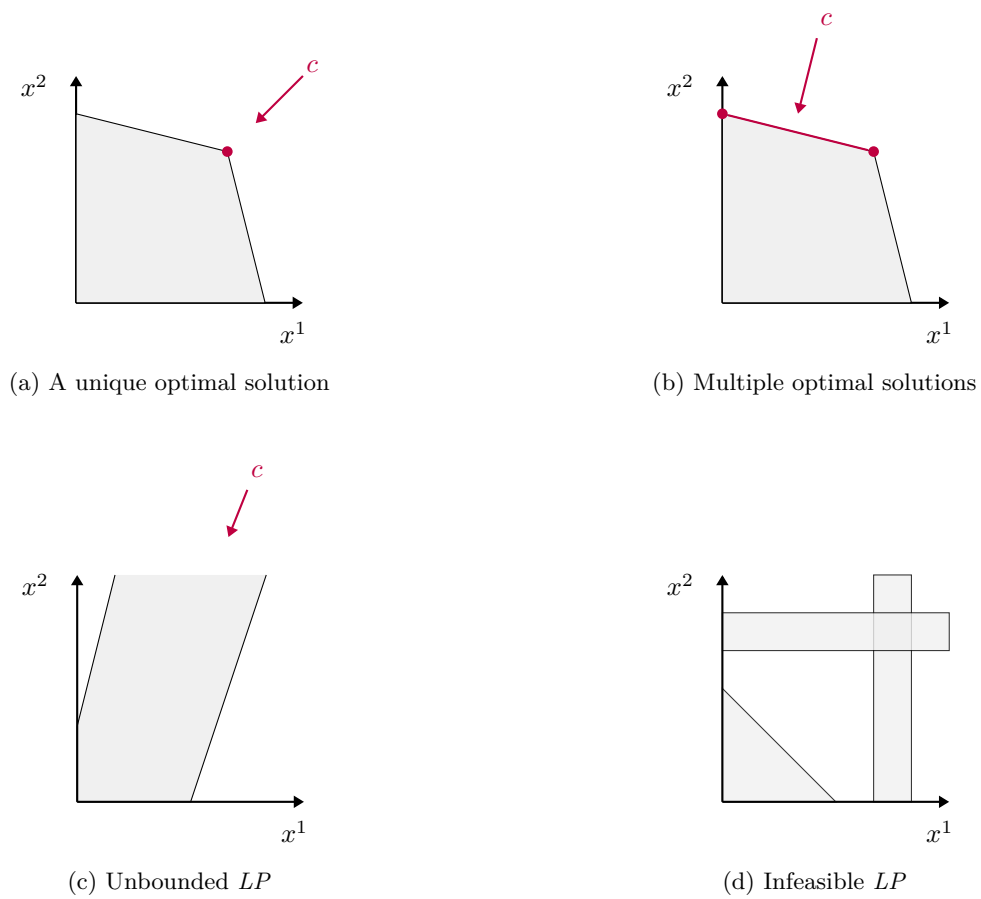
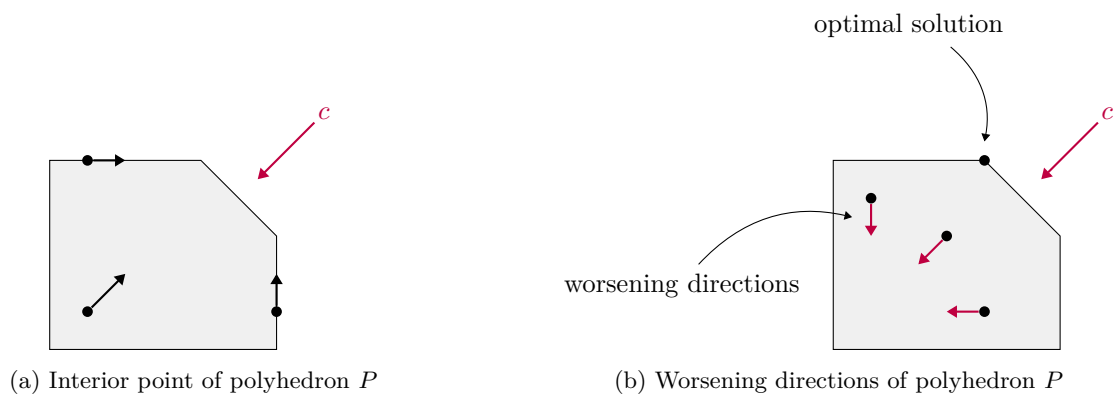


Figure 16: Four types of LP

3.4.4 Basic feasible solutions and polyhedra vertices

Due to the fundamental theorem of Linear Programming, to solve any LP problem it suffices to consider the (*finitely many*) vertices of the polyhedron P of feasible solutions. Since the geometrical definition of vertex cannot be exploited algorithmically, an algebraic definition is needed.

A vertex corresponds to the intersection of the hyperplanes associated to n inequalities. If a polyhedra is expressed in standard form

$$P = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$$

it is possible to transform it into a inequality

$$P = \{x \in \mathbb{R}^n \mid Ax \leq b, x \geq 0\}$$

and later transform it into standard form

$$P' = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$$

where P' is the polyhedron of feasible solutions of the original problem.

Finally, by renaming

$$A := [A \mid I] \quad x := (x^T \mid s^T)$$

the system of equation is represented in matrix form, where A has m rows.

Property 3.1. For any polyhedron $P = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$, where A has m rows:

- the **facets** (edges in \mathbb{R}^2) are obtained by setting one variable to 0
- the **vertices** are obtained by settings $n - m$ variables to 0

3.4.5 Algebraic characterization of vertices

Consider any polyhedron $P = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$, in standard form.

Assumption

$A \in \mathbb{R}^{m \times n}$ is such that $m \leq n$ of rank m (*i.e.* A is full rank). This is equivalent to assume that there are no redundant constraints.

Solutions

- If $m = n$, there is a unique solution of $Ax = b$ ($x = A^{-1}b$)
- If $m < n$, there are ∞^{m-n} solutions of $Ax = b$
 - the system has $n - m$ **degrees of freedom**
 - by fixing the degrees of freedom to 0, a **vertex** is obtained

Definition 3.4. The **basis** of matrix A is a subset of m columns of A that are linearly independent and form an $m \times m$ non singular matrix B .

$$A = \left[\overbrace{B}^m \mid \overbrace{N}^{m-n} \right]$$

3.4.6 Basic solutions

Let $x^T = \left[\overbrace{x_B^T}^{m \text{ components}} \mid \overbrace{x_N^T}^{m-n \text{ components}} \right]$. Then any system $Ax = b$ can be written as $Bx_B + Nx_N = b$, and for any set of values of x_N , if B is not singular, then $x_B = B^{-1}(b - Nx_N)$.

Definitions

- A **basic solution** is a solution obtained by setting $x_N = 0$ and, consequently, $x_B = B^{-1}b$
- A **feasible solution** is any vector $x \geq 0$ such that $Ax = b$

- A feasible solution with $x_B \geq 0$ is a **basic feasible solution** with basis B
 - $\forall j \notin B, x_j = 0$ - all the non zero variables are in B
 - a **basic feasible solution** has n linearly independent constants
 - the variables in x_B are the basic variables and those in x_N are non basic variables
 - by construction, (x_B^T, x_N^T) satisfy $Ax = b$
 - a basic feasible solution is referred to as *BFS* for brevity
- A basic feasible solution is **degenerate** if it contains at least one basic variable with value 0

Theorem 3.2 (Basic feasible solutions). $x \in \mathbb{R}^n$ is a basic feasible solution if and only if x is a vertex of the polyhedron $P = \{x \in \mathbb{R}^n \mid Ax = b, x \geq 0\}$

Property 3.2 (Number of solutions). At most, there exists one basic feasible solution for each choice of the $n - m$ non basic variables out of the n variables:

$$\# \text{ basic feasible solutions} \leq \binom{n}{n-m} = \frac{n!}{(n-m)!(n-(n-m))!} = \binom{n}{m}$$

3.5 Simplex method

The Simplex method provides a systematic way to find the optimal solution of a *LP* problem. Given an *LP* problem in standard form:

$$\begin{aligned} \min \quad & z = c^T x \\ \text{s.t.} \quad & Ax = b \\ & x \geq 0 \end{aligned}$$

it examines a sequence of basic feasible solutions with non increasing objective function values, until an optimal solution is reached or the problem is found to be unbounded. At each iteration, the basic feasible solution is replaced by a new one that is closer to the optimal solution.

In other, *simpler*, words, it generates a path (*a sequence of adjacent vertices*) along the edges of the polyhedron P that leads to the optimal solution.

The simplex method is articulated in 3 steps:

1. Find an **initial vertex** or establish that the *LP* is **unbounded**
2. Determine whether the current vertex is **optimal**
3. Move from a current vertex to a **better adjacent vertex** (*in terms of objective function value*) or establish that the *LP* is unbounded

3.5.1 Optimality test

Given a *LP* $\min \{c^T x \mid Ax = b, x \geq 0\}$ and a feasible basis B of A , $Ax = b$ can be written as:

$$B_{x_B} + N_{x_N} = b, \quad B^{-1}b \geq 0$$

where x_B and x_N are the basic and non basic variables, respectively. Then, a **basic feasible solution** is a value of x such that $x_N = 0$, and consequently $x_B = B^{-1}b$.

By substitution, the objective function can be expressed in terms of only the non basic variables:

$$z = c^T x = c_B^T B^{-1}b + (c_N^T - c_B^T B^{-1}N) x_N$$

- The cost of $x = (x_B, x_N)$ is $z_0 = c_B^T B^{-1}b$ where $x_B = B^{-1}b, x_N = 0$
- The reduced costs of the non basic variables x_N are denoted by $\bar{c}_N^T := c_N^T - c_B^T B^{-1}N$

- The vector of reduced costs with respect to the basis B is

$$\bar{c}^T = c^T - c_B^T B^{-1} A = \overbrace{[c_B^T - c_B^T B^{-1} B]}^{0^T}, \overbrace{[c_N^T - c_B^T B^{-1} N]}^{\bar{c}_N^T}$$

- The reduced costs are defined also for basic variables, but $\bar{c}_B = 0$
- \bar{c}_j represents the change in the objective function value if non basic variable x_j is increased from 0 to 1 while keeping all other non basic variables to 0
- The solution value changes by $\Delta z = \theta \bar{c}_j$
- The solution value changes by $\Delta z = \theta \bar{c}_j$

Proposition 3.1. If $\bar{c}_N \geq 0$, then the basic feasible solution (x_B^T, x_N^T) , where $x_B = B^{-1}b \geq 0, x_N = 0$, of cost $z_0 = c_B^T B^{-1}b$ is a global optimum.

Proof. $\bar{c}^T \geq 0^T$ implies that

$$c^T x = c_B^T B^{-1}b + c_N^T x_N \geq c_B^T B^{-1}b \quad \forall x \geq 0, Ax = b$$

□

Remarks

- For maximization problems, the condition is $\bar{c}_N \leq 0$.
- This optimality condition is sufficient but generally not necessary.

3.5.2 General case

Given a basis B , the system $Ax = b$ can be written as:

$$\sum_{j=1}^m a_{ij}x_j = b_i \quad \text{for } i = 1, \dots, m$$

and then expressed in canonical form

$$x_B + B^{-1}N x_N = B^{-1}b \Leftrightarrow x_B + \bar{N} x_N = \bar{b}$$

which emphasizes the basic feasible solution $(x_B, x_N) = (B^{-1}b, 0)$. This amounts to pre-multiply the system by B^{-1}

$$\underbrace{B^{-1}B x_B}_{=I} + \underbrace{B^{-1}N x_N}_{=\bar{N}} = \underbrace{B^{-1}b}_{=\bar{b}}$$

In the canonical form

$$x_{B_i} + \sum_{j=1}^{n-m} \bar{a}_{ij} x_{N_j} = \bar{b}_i \quad i = 1, \dots, m \quad (I_{x_B} + \bar{N}_{x_N} = \bar{b})$$

basic variables are expressed in terms of non basic variables:

$$x_B = \bar{b} - \bar{N}_{x_N}$$

If the value of value of a non basic x_s is increased from 0 while all the other non basic variables are kept to 0, the system is modified as follows:

$$x_{B_i} + \bar{a}_{is} x_s = \bar{b}_i \Leftrightarrow x_{B_i} = \bar{b}_i - \bar{a}_{is} x_s, \quad i = 1, \dots, m$$

In order to guarantee $x_{B_i} \geq 0 \forall i$, the value of x_s must be chosen such that

$$\bar{b}_i - \bar{a}_{is}x_s \geq 0 \Rightarrow x_s \leq \frac{\bar{b}_i}{\bar{a}_{is}} \quad \bar{a}_{is} \geq 0$$

Then the value of x_s can be increased up to

$$\theta^* = \min_{i=1}^m \left\{ \frac{\bar{b}_i}{\bar{a}_{is}} \mid \bar{a}_{is} \geq 0 \right\}$$

while the value of x_r , where

$$r = \arg \min_{i=1}^m \left\{ \frac{\bar{b}_i}{\bar{a}_{is}} \mid \bar{a}_{is} \geq 0 \right\}$$

decreases to 0 and exists the basis.

If $\bar{a}_{is} \leq 0 \forall i$, there's no limit to the increase of x_s .

3.5.3 Changes of basis

Assumption: this Section refers to the case where the LP is a minimization problem.

Let B be a feasible basis and $x_s \in x_N$ a non basic variable with reduced cost $\bar{c}_s \leq 0$. In order to change the basis:

- x_s is increased as much as possible (x_s “enters the basis”) while keeping the other non basic variables equal to 0
- The basic variable x_r (in x_B) such that $x_r \geq 0$ imposes the tightest upper bound θ^* on the increase of x_s (x_r “leaves the basis”)
- If $\theta^* \geq 0$, the new basic feasible solution has a better objective function value than the previous one
- The new basis differs from the previous one by a single column (*adjacent vertices*)
-

To transform the canonical form of the current basic feasible solution into the canonical form of the new basic feasible solution, the following formula is applied

$$B^{-1}Bx_B + B^{-1}Nx_N = B^{-1}b$$

The value B^{-1} of the new basis B can be obtained by applying to the inverse of the previous basis a single elementary row operation, called **pivot**; it will be explained in the next Paragraph.

3.5.4 Pivoting operation

The pivot operation is a single elementary row operation that transforms the inverse of the previous basis into the inverse of the new basis; it's the same operation used in the Gaussian elimination method to solve systems of linear equations.

Given $Ax = b$:

1. Select a coefficient $\bar{a}_{rs} \neq 0$, the **pivot**
2. Divide the r -th row by \bar{a}_{rs}
3. For each row $i \neq r$ with $\bar{a}_{rs} \neq 0$, subtract the resulting r -th row multiplied by \bar{a}_{is}

This move does not affect the feasible solutions. Only a finite number of pivots exists for each basis.

An example of pivot operation is shown in Figure 17.

		s					
		↓					
			1	1	1	0	6
r →			2	1	0	1	8

(a) Initial system

		s					
		↓					
			0	$1/2$	1	$-1/2$	3
r →			1	$1/2$	0	$1/2$	4

(b) Pivot operation

Figure 17: Pivot operation example

3.5.5 Moving to an adjacent vertex

Goals

- improve the objective function value
- preserve the feasibility

1. Which non basic variable enters the basis?

- any one with reduced cost $\bar{c}_j < 0$
- the one yielding the maximum Δz with respect to the current basic feasible solution ($z = c_B^T B^{-1}b$)
- **Bland's rule:** $s = \min\{j \mid \bar{c}_j < 0\}$
 → $\bar{c}_j > 0$ for maximization problems
- Which basic variable leaves the basis?
 - **min ratio text:** index i with smallest positive ratio $x_{B_i}/\bar{a}_{is} = \theta^*$ among those with $\bar{a}_{is} > 0$
 - **Bland's rule:** $r = \min\{i \mid \bar{b}_i/\bar{a}_{is} = \theta^* \bar{a}_{is} > 0\}$
 - also randomly

Property 3.3 (Unboundedness). If the objective function is unbounded, the algorithm will never stop: if $\exists \bar{c}_j < 0$ with $\bar{a}_{ij} \leq 0 \forall i$ no element of the j -th column can play the role of the pivot. This condition is verified if problem is **unbounded**.

3.5.6 Tableau representation

Tableau is a matrix representation of the LP problem. Let

$$z = c^T x$$

$$Ax = b$$

with implicit non negativity constraints be the LP problem.

The initial tableau is the matrix represented in Figure 18a. The first column contains the right hand side of the objective function and the right hand side vector.

Consider a basis B and a partition $A = [B \ N]$, with $0 = c^T x - z$. The corresponding tableau is the matrix represented in Figure 18b.

By pivoting operations (*or pre multiplication by B^{-1}*), the tableau is in the canonical form with respect to B . The said tableau is the matrix represented in Figure 18c.

3.5.7 The simplex algorithm

In Code 8 the simplex algorithm is implemented.

```
B[1], ..., B[m] := initial basis
construct the initial tableau A in canonical form with respect to B
```

	x_1	\cdots	x_n	
0	c^T			← <i>objective function</i>
b	A			← <i>m rows</i>

(a) Initial tableau

	x_1	\cdots	x_m	x_{m+1}	\cdots	x_n	z
0	c_B^T			c_N^T			-1
b	B			N			0 \vdots 0

(b) Tableau with respect to a basis

	x_1	\cdots	x_m	x_{m+1}	\cdots	x_n	
$-z_0$	0	\cdots	0	\bar{c}_N^T			
\bar{b}	I			\bar{N}			

(c) Canonical tableau

Figure 18: Tableau representation

```

unbounded := false
optimal := false
while optimal = false and unbounded = false do
  if a[0, j]  $\forall j = 1, \dots, m$  then
    optimal := true // for LP with min
  else
    select a non basic variable x_s with a[0, s] < 0 // negative reduced cost
    if a[i, s]  $\leq 0 \forall i = 1, \dots, m$  then
      unbounded := true
    else
      r := argmin {a[i, 0] / a[i, s]  $\forall i = 1, \dots, m$  with a[i, s] > 0}
      pivot(r, s) // update the tableau
      B[r] := s
    end
  end
end

```

Code 8: The simplex algorithm

3.5.8 Degenerate basic feasible solutions and convergence

As already announced in Section 3.4.6, a basic feasible solution x is degenerate if it contains at least one basic variable $x_j = 0$. A solution x with more than $n - m$ zeroes corresponds to several distinct bases.

More than n constraints (*the m of $Ax = b$ and more than $n - m$ among the n of $x \geq 0$*) on the same vertex are satisfied with equality. In the presence of degenerate basic feasible solutions (*BFS*), a basis change may not decrease the objective function value: if the current *BFS* is degenerate, the only admissible value of θ^* is 0 and the new *BFS* is the same as the old one.

Note that a degenerate *BFS* can arise from a non degenerate one: even if $\theta^* > 0$, several basic variables may go to 0 when x_s is increased to θ^* . It's possible to cycle through a sequence of degenerate basis associated to

the same vertex.

3.5.8.1 Anti cycling rule

Several anti cycling rules have been proposed for the choice of the variables that enter and exit the bases (*indices r and s in Algorithm 8*).

The simplest one is the **Bland's rule**: among all candidate variables for x_s and x_r , the one with the smallest index is chosen.

Property 3.4. The simplex algorithm with Bland's rule terminates after less than $\binom{n}{m}$ iterations.

3.5.9 Two phase simplex algorithm

The two phase simplex algorithm is a modification of the simplex algorithm that allows to solve the *LP* problem with equality constraints.

Phase 1: Determine an initial basic feasible solution

- given the problem with equality constraints

$$\begin{aligned} z &= c^T x \\ Ax &= b, b \geq 0 \\ x &\geq 0 \end{aligned}$$

- an auxiliary *LP* with artificial variables is constructed

$$\begin{aligned} \min \quad & v = \sum_{i=1}^m y_i \\ \text{s.t.} \quad & Ax + ly = b \\ & x \geq 0, y \geq 0 \end{aligned}$$

Phase 2: Solve the auxiliary problem

- if $v^* > 0$, the problem is infeasible
- if $v^* = 0$, $y^* = 0$ and x^* is a basic feasible solution of the original problem
 - if y_i is non basic $\forall i$, with $1 \leq i \leq m$, the corresponding columns are deleted and a tableau in canonical form is obtained; the row of z must be determined via substitution
 - if there is a basic y_i (*the basic feasible solution is degenerate*), then a pivot operation is performed on the row of y_i to exchange the basic variable with a non basic one

3.6 Network flows

Network flows problems involve the distribution of a given *product* (*such as water, gas, data*) from a set of *sources* to a set of *users* so as to optimize a given objective function (*e.g. minimize the total cost of the distribution*).

It has many indirect applications, such as:

- **Telecommunication**
- **Transportation**
- **Logistics**

Definition 3.5 (network flow). A network is a directed and connected graph $G = (V, A)$ with a source $s \in V$ and a sink $t \in V$, with $s \neq t$, and a capacity $k_{ij} \geq 0$ for each arc $(i, j) \in A$.

- A **feasible flow** x from s to t is a vector $x \in \mathbb{R}^m$ with a component x_{ij} for each arc $(i, j) \in A$ satisfying the capacity constraint

$$0 \leq x_{ij} \leq k_{ij}, \quad \forall (i, j) \in A$$

and the flow balance constraint at each intermediate node $u \in \{V \setminus \{s, t\}\}$:

$$\sum_{(i,u) \in \delta^-(u)} x_{iu} = \sum_{(u,j) \in \delta^+(u)} x_{uj} \quad \forall u \in N \setminus \{s, t\}$$

- The value of **flow** x is $\phi = \sum_{(s,j) \in \delta^+(s)} x_{sj}$
- Given a network and a feasible flow x , an arc $(i, j) \in A$ is **saturated** if $x_{ij} = k_{ij}$ and **empty** if $x_{ij} = 0$

Definition 3.6 (flow problem definition). Given a network $G = (V, A)$ with an integer capacity k_{ij} for each arc $(i, j) \in A$, find a feasible flow x from s to t with maximum value.

If multiple sources or sink are present while only one product is considered, dummy nodes s^* and t^* can be added.

3.6.1 Linear programming model

The linear programming model of the network flow is defined as:

$$\begin{aligned} \max \quad & \phi \\ \text{s.t.} \quad & \end{aligned} \tag{3.6}$$